Peer-To-Peer in Botnets

Moritz Marc Beller

Fakultät für Informatik,
Technische Universität München
`{beller}@in.tum.de`

Abstract. In this literature review paper, we describe a relatively novel means of communication in botnets, namely *peer-to-peer* (P2P) communication. While many users are familiar with P2P file sharing, the term has attracted more attention lately as malicious botnets such as Storm[DFNM] have emerged, causing quantifiable real-world damage.
We explain P2P botnets and compare them to conventional ones. We describe possible countermeasures against P2P botnets and give an outlook into the future of botnets.

1 Introduction

Cyber-attacks are a growing threat to the security of the world. As a result, in June of 2011, the German government founded a "national defence centre against cyber crime".[fSidIB11] This step became necessary as new forms of computer worms are no longer restricted to harming only computers and the data stored on them, as the example of "Stuxnet" demonstrates:

In 2010, a new computer worm called "Stuxnet" was discovered. It is suspected to have damaged Iran's uranium centrifuges, giving it the title "the first cyber weapon"[Ben11]. To spread and update, Stuxnet used Internet and Intranet structures to build its own *botnet*.[FMC] Although Stuxnet didn't try to build an armada of infected computers like most other bots do, Stuxnet is in alignment with a series of new botnets using *peer-to-peer* communication instead of a centralised server.

This review paper focuses on these peer-to-peer botnets and the growing threat emerging from them. Having read this article, the reader will be well-informed about current and future techniques in peer-to-peer botnets.

Outline

We start by introducing common botnet terminology. We then give a classification of P2P botnets and distinguish different architectures of botnets. In the following sections, we describe the typical lifetime of a P2P botnet, the command and control structure in P2P botnets, comparing conventional botnets to P2P botnets. In the remainder of this review paper we described possible countermeasures against P2P botnets and draw conclusion on what types of botnets to expect in the future.

2 Definitions

A computer able of executing commands remotely triggered by a malicious attacker and possibly against the will of the owner is called a *bot* or *zombie*. The term can also generally refer to the program executed, meaning not a special instance of an infested computer, but the software or idea behind the software per se (as in "the Storm bot").

A *botnet* is a group of bots forming a common network structure.[SK07] In most recent papers on the subject ([WWAZ09], [ARZMT06]), the term botnet is defined

as purely negative, i.e. a network performing destructive aims such as denial-of-service attacks[1], sending spam or hosting a phishing website[SI07]. Other common aims include providing the aggregated CPU resources of the botnet, stealing users' credentials [Bor] or doing click fraud.[2] on affiliate networks[New11]

In the following, we use a bias-free definition of botnet as per our understanding technology is generally ethics-free. Additionally, there are many examples where botnets are used in a non-destructive way (e.g. [oC11]), or even to destroy existing malicious botnets.

The controller of the botnet is referred to as the *botmaster*. He doesn't necessarily have to be the founder of the botnet (cf. section 3.2).

The expression *bot candidates* specifies the set of computers which are target to becoming a bot themselves.

P2P, short for *Peer-to-Peer*, being a technology buzz word of the internet in the late 1990s with the upcoming file sharing services like Napster[Inc11], has attracted less attention in recent years. P2P defines an unstructured information network amongst equals — so-called peers. Two or more peers can spontaneously exchange information without a central instance. According to [SFS05] "P2P networks promise improved scalability, lower cost of ownership, self-organised and decentralised coordination of previously underused or limited resources, greater fault tolerance, and better support for building ad hoc networks." These properties coupled with the fact that files exchanged in P2P networks are prone to malware, trojans and viruses make P2P networks a most-attractive base for building botnets. Well-known P2P networks include the Napster, Gnutella, Overnet and Torrent network.

A *P2P bot* then is a bot that uses a P2P protocol as a means of communication with other bots.

The so-called *C&C*, command and control structure, specifies the way and protocols in which the botmaster and the bots communicate with each other. It is the central property of any botnet. Common protocols for C&C include IRC, HTTP, FTP and P2P.[Bor]

IRC — internet relay chat — is a "teleconferencing system"[irc], typically used for text chatting in channels joined by a large number of participants. While its protocol is relatively easy to implement, it provides a lot of features. It has thus become the one of the most widely-used protocols for C&C in conventional botnets.

The process of *bootstrapping* generally describes starting a more complex system out of a simple system. In regard to botnets, the term usually means loading of the bot code (often injected into the original file sharing program) and establishing a connection to other bots.[WWAZ09]

3   An overview over P2P botnets

This section shall give a short overview of P2P botnets: How they came into being, how one can classify them according to their distribution strategy and what the typical life of a botnet looks like.

3.1   A brief history of botnets

The origins for the term "bot" go back to so-called IRC bots, short for IRC robots. An IRC bot is a program that handles specific tasks in the IRC automatically, so that the administrator does not need to do those routine jobs.

---

[1] Abbreviated: DoS attacks, i.e. generating so much traffic the target cannot function normally any more

[2] Click fraud is the process of generating clicks on web banners without an actual user seeing or clicking the advertisement, for the sake of making money.

The first IRC bot ever to be created was named Eggdrop. Its origins go back to the year 1993. However, in April 1998 a deviant called GT-Bot appeared and formed the first malicious botnet, using IRC's C&C structures. Four years later, in 2002, Slapper was the first botnet to make use of P2P for C&C.[LJZ]

There have been many more P2P botnets created since, most notably perhaps the Phatbot and Storm botnets. Phatbot (also called Agobot, Gaobot) became renowned for its large set of built-in attacks ranging from password sniffing to DoS.[CJM05] The storm botnet, which exists in several different versions, was described "as one of the most sophisticated botnet[s] active today"[DFNM]. Storm is based on the Overnet network, a P2P network established by programs such as eMule and eDonkey.

A criminal conviction because of launching a botnet is relatively seldom, since botmasters try to hide their identities (see "stepping stones" in figure 1). Yet, in 2007, John Schiefer was sentenced to four years in prison. He built a botnet with up to 250,000 zombies, collecting passwords and bank credentials from the bots.[Hru11]

3.2 Classification of P2P networks

There are three types of P2P networks: *parasite, leeching* and *bot-only*.[WWAZ09]

Parasite and leeching bots infiltrate existing P2P networks, while "bot-only" networks are designed with new network protocols.

Parasite botnets recruit new bots only from the set of existing P2P participants; they try to infect system inside the P2P network and make them become bots. Due to the often illegal content distributed in file sharing networks, they are a perfect culture medium for viruses, malware and worms. It is thus convenient for an attacker to spread a highly-demanded file (e.g. some cracked computer game, software or porn) containing the injection code sequences for his bot. This code is then injected into the file sharing client. Vulnerable hosts in the network are infected this way. On the downside, this means that the spread of the bot is limited to the size of the P2P network.

In contrast, leeching bots not only try to infiltrate systems which are already part of the P2P network, but also systems outside of the P2P network. Naturally, they are bigger in size as they have to deliver the P2P client, too. This might be more difficult to achieve as it means that systems must unwillingly take part in the network. Often, firewalls and port-forwarding are not configured for use with the P2P network on these systems, reducing the performance of the botnet: Since the owner of the computer usually doesn't even know he is participating in a P2P network, he has no intention of opening the ports in his firewall. Leeching bots can spread through any possible measure: File sharing, downloads on websites, email attachments or instant messaging, to name a few.

There are good reasons for either strategy — relying on existing networks or creating a new one: Using an extant P2P network as a base — like parasite and leeching bots do — unburdens the botmaster from setting up and building a botnet infrastructure. It profits from the established P2P network, making use of filtering, error-correction and encryption as far as the chosen network has support for it. On the other hand, features are limited to the existing P2P protocol. A specifically-built P2P bot-only network is naturally more tailored towards its purpose. Due to the bot-exclusive memberships, it might be easier to shutdown the botnet for an attacker as all participants can be considered bots and there is no risk of accidentally shutting down an innocent member. Bot-only networks are described in detail in section 5.2.

3.3 Lifetime of P2P botnets

Wang et al.[WWAZ09] differentiate three stages of P2P botnets:

- Recruiting bot members (infecting others)
- Forming the botnet (construction phase)
- Standing by for instruction
  This is the actual "operational" phase of the botnet. Bots are awaiting instructions from their master. Instructions can either be actual commands or performing updates. In this phase, the chosen C&C structure is essential.

It should be noted that these phases are not strictly exclusive, e.g. during the third phase (standing by for instruction), building of the botnet may well continue. In fact, this is an inherent property of any P2P network: Constant transformation of the network. It is only until a critical mass of bots has proceeded past phase one and two, that the botnet can be called operational.

When speaking about P2P networks, a factor that must be taken into consideration is the *churn* of the network. By this name, we describe the collective effects that the independent connects and disconnects of thousands of peers can have. The term churn also refers to the different session lengths (most P2P users typically have their PC running for several hours, while there are some which stay constantly online). In short, churn stands for the dynamics in a P2P system. Despite its great significance, churn has not been discussed in the literature a lot. [SR06]

4 Command and Control structures in botnets

In this section, we give an overview of the C&C typically used in P2P botnets and compare it to the C&C in centralised architectures. The differences in C&C are the most notable differences between the two, apart from the architecture itself (cf. section 5). There exist three fundamentally different forms of C&C in botnets: Centralised C&C, P2P C&C and unstructured C&C. After clarifying the differences between push and pull communication, we explain the three approaches.

4.1 Push/Pull mechanism

Commands can be distributed in two ways: Either via a push or a pull mechanism. "Pulling", being the more trivial of the two approaches, is the process of a client actively asking a server whether there's new instructions for it. To be up-to-date, this has to be performed periodically, increasing network load. Push on the other hand is technically more advanced, as commands from the server will be automatically "forwarded" to the clients — the difference between push and pull in a bot is much like IMAP IDLE compared to POP3 for mail: With POP3, the user client has to periodically ask the server whether there are new messages available (pull), whereas with IMAP, the client gets a new message pushed by the server. This has the advantage of reducing the network traffic, but it requires that the server can open a connection to the clients (cf. section 5.3).

4.2 C&C in centralised networks

C&C in centralised networks is more straight-forward than in P2P networks. The central servers receive their instructions directly from the botmaster. Only he has access to those resources, so there is no problem of authentication as in P2P networks (cf. section 6.1). Subsequently, the servers forward their commands to all their clients, or the clients ask their servers for new orders, depending on whether push or pull communication is used (cf. section 4.1). When information from the bots is gathered, this is collected on the central servers, ready to be downloaded by the botmaster. When using IRC as a protocol for C&C (another alternative would

be HTTP, [LJZ]), the central server has an IRC server running on a specific port which the bots connect to. The bot must thus contain a suitable IRC client or IRC protocol implementation. The exact procedure differs from botnet to botnet, but is usually similar to the following: On the server, IRC channels are setup which the bots join. Often, all the bots are set to listen-only. The botmaster is the controller of the channel. He can send commands through private messaging to specific bots, or broadcast commands by "chatting" to the whole channel. Additional services like IRC file transfers can be used to update the bot, or to receive files from the bots.

4.3   C&C in structured P2P networks: File indices

Almost all P2P file sharing systems have so-called *indices*, in which they map a desired content to a location where it can be found (IP address and port). An index can thus be considered as a "lookup table" for finding the information where a user can retrieve a certain content he searches for. In terms of botnets, this usually resembles the place where a bot can find its orders. Therefore, indices are a vital part of any such P2P system employing them.

File indices are typically implemented using a hashed map. For botnets, file indices are usually used not to publish the location of a file, but to distribute commands. One record in such a map thus has the format $< k, m >$, where k is the key under which to store the command m.

In order to find a search result for a query q, in structured P2P networks such as Overnet, the following algorithm is used: q is directed to nodes whose IDs are closer to the queried hash key. If it finally finds a peer which can offer the searched k, a direct connection between the two is established (peer-to-peer). Otherwise, if no k can be found, the search will eventually hit the timeout and fail. From a botmaster's perspective, it makes sense to associate one command with a number of different k. In Storm's botnet for example the bots query for freshly generated keys every day.[WWAZ09] The key to be queried by a single bot is calculated as a function of date (uniform for all bots) and a random number $\in [0; 31]$ (varies between every bot) locally on every bot. Therefore, the botmaster has to publish every command under 32 different keys for all bots to be able to receive the command. This way, the command is sent to other sub-networks, distributing the command nicely within the botnet.

Depending on the architecture of the P2P net (cf. section 5), the index might be stored in a central position (Napster-like file sharing systems, see section 5.2 on page napsterlike), or be distributed over a larger number of nodes. In the latter case, usually only a part of the file index is available on a peer. If a peer doesn't have the key k a client is searching for in its hash map, it forwards this search to its neighbours, so that in the end, the client will receive a pair of $< k, m >$ for any given k, given the network's churn is healthy and there exists a k at all.

Instead of providing information on where to find a certain content as in file sharing networks, botmasters insert the actual command to be executed by the bots into the record. In order to get commands issued by botmasters, bots periodically initiate queries for a set of k, and nodes which preserve the corresponding records will return query hits with the commands encoded. Thus, using file indices implies using pull mechanism. [LNR06]

4.4   C&C in unstructured P2P networks

In contrast to the structured P2P C&C, in unstructured C&C a bot won't ever actively contact other bots. Instead, it listens to incoming connections. [LJZ] Thus,

instead of searching a path to the nearest node with a similar hash ID, in unstructured P2P, the network gets "flooded" with query messages, generating more overhead traffic.

When commands are pushed in P2P networks, this implies a more complex infrastructure. [WWAZ09] formulate two problems to be dealt with:

– **To which peers should a received command be forwarded?**
One solution would be to distribute a command to a bot's neighbouring peers. However, this might make command distribution slow and incomplete, since many bots have a small number of neighbours (cf. figure 3 on page 10). Additionally, in parasite and leeching botnets not all neighbours are running the bot, thereby not being able to receive the commands.
Another solution to the problem is for the bots to search for some predefined Keys k. All bots are programmed in such a way that they answer to have that Key available. Then commands are distributed to all the peers that answered the search for the k. However, this might make the botnet target-able to attacks, as one only has to know the right k to find all bots. [WWAZ09] call those k figuratively "watchwords".
– **How should the commands be forwarded?**
There are two possible ways of forwarding a command: Using *in-band* (normal P2P traffic) or *out-of-band* (non-P2P traffic) communication. This question is only relevant to leeching and parasite botnets.
In-band messages are the safest choice in terms of robustness of the botnet, but only work if a bot targets its neighbouring peers[WWAZ09]. By using in-band traffic, it is hard for attackers to differentiate between "normal" P2P traffic and the botnet's C&C. In-band messages have the downfall that a bot cannot communicate with more remote peers, but only with those peers that it would normally exchange information with. This might cause command distribution to be slow and incomplete.
Out-Of-Band messages do not share this restriction, as any bot can simply open a connection with any other bot in the network. However, it is easy to detect and trace out-of-band messages, since they stand out from the usual P2P traffic. How can a bot know about another bot which is not in its neighbouring peer list at all? It could have been found through a search for certain watchwords (see above). Then, to contact the other bot for exchange of commands, a special request for connection has to be made. Exactly this step could disguise the two bots to an attacker of the botnet.

5 Architectures of Botnets

In reality, there exist two principally different architectures of botnets. A new architecture has been proposed by Wang et al.[WSZ10], which combines the advantages of both. The following nomenclature is extracted from [SI07]:

5.1 Centralised Architecture

Historically the oldest form of botnets, centralised architectures are built up in such a way that there is one central spot which broadcasts messages between the connected bots and the botmaster. It's common to have more than one central server[WSZ10], but even with several servers, the architecture still stays centralised. For if you shutdown this central point, the network is inoperable. This resembles the biggest weakness of centralised architectures: As soon as you are able to cut the server off the net, the network is dead. On the other hand, latency becomes minimal,
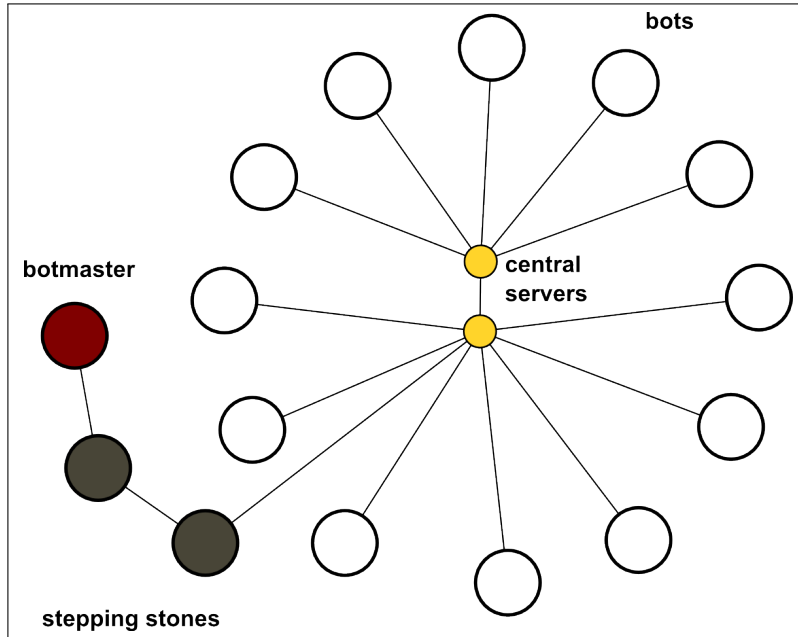
Fig. 1. Graph depicting connections in a centralised network. Note how all bots only have a connection to the central point.

as the routing distance for one package needed to reach each node in the botnet is minimal (only one transition is needed). Bandwidth, however, is generally limited by the server's resources, making it hard to receive or transmit big chunks of data. Furthermore, it holds that all the routes in the botnet have the same length. This is something which is fundamentally more complex in decentralised networks and can be a great simplification of command distribution and monitoring the network.

Due to their nature, centralised architectures are usually implemented with an IRC C&C or similar[CJM05]. The central server is normally not owned by the bot-master. This would make detecting his identity easy. In many countries, launching an "evil-minded botnet" is a serious crime. Instead, hacked or public IRC servers are used as the central C&C node. A connection from the attacker's computer to the central server is often obfuscated by many in-between relays, tunnels and encryption. In figure 1 this is shown as the "stepping stones" which shall hide an attacker's identity.

5.2 Decentralised Architecture

Decentralised architectures do not rely on the special role of one central server. Instead, they are built upon the principal of equality, namely that the "peer nodes (both client and server) are all equal"[SI07]. The topology of the network is far more complex than in centralised architectures, forming a mesh as shown in figure 2. It is thus more difficult for a bot to join the botnet. Extensive bootstrapping is required, as the bot has to figure out an already-participating peer to connect to in the beginning. Once inside the net, information about other peers is exchanged between nodes. There are two approaches for bootstrapping[WWAZ09]:

- A list of peers likely to be online is hard-coded into the client. This list can later be updated
- A shared web cache on the internet stores information about peers. The address of which is hard-coded.

As can be seen, Bootstrapping is a critical and vulnerable point in any P2P botnet. Considerable efforts by botmaster have been made to circumvent the need to bootstrap[WSZ10]. This is further discussed in section 6.

Once inside the net, information about other peers is exchanged between nodes.

Distributing commands and data in such a network is complicated, as it has to be assured that the message reaches all clients. As a general rule of thumb, the better inter-connected the nodes are, the higher the probability for a message to reach all recipients.

A decentralised network has the advantage of having the accumulated resources and bandwidth of all the peers in the network available. However, latency might be bad, as routing through the network is not trivial (cf. figure 2). P2P networks are generally considered to be harder to disable (cf. section 6 on page 13). Another advantage might be that the botmaster does not necessarily need to run server of his own — decreasing infrastructure costs in comparison with centralised architectures.

It is to be discussed whether P2P networks with a centralised server architecture for certain services like file-indexing — we refer to them as "Napster-like" botnets — fall into the "decentralised" category. Principally, the connection graph differs a lot from centralised networks, but centralised and Napster-like networks share the same weaknesses, as could be seen when Napster was shut down in 2001[Wik11].[3] Dittrich et al. [DD07] would consider Napster-like botnets a hybrid architecture, whereas Steggink et al. [SI07] classify it as decentralised. In the following, we stick with Steggink's definition, as there seems to be a broader consensus in the literature towards their nomenclature (cf. [WSZ10]).

5.3 Hybrid Architecture

Hybrid architectures are botnet architectures driven by scientific development and up to this date only exist in theory. They have been described as "the advanced P2P botnet"[WSZ10] and the "super botnet"[VAJ07]. The approach is to study current botnets, analyse their weaknesses and propose a better solution. This anticipates how botmasters could improve their botnets in the future. This way, even today, we know what future botnets could look like and how to better defend against them.

The proposed new hybrid P2P botnets do not have a pre-set communication architecture, following the strict P2P definition. Their network connectivity is solely determined by the *peer list* in each bot.

Only machines with static IPs appear as bots in the peer list, so-called *servents*. [WSZ10] This way, it is guaranteed that the distributed peer lists are maximally deadlink-free. This is a specialisation of the "all peers are equal" contract in P2P: Some bots — the servents — have special obligations described in the following. The clients are then typically bots behind firewalls, machines with private or dynamic IPs.

Network construction phase Infection is done no different than in conventional botnets (cf. section 3.2). The basic construction procedure has two mechanisms:

- **New Infection:** "Bot A passes its peer list to a vulnerable host B when compromising it. If A is a servent bot, B adds A into its peer list (by randomly replacing one entry if its peer list is full). If A knows that B is a servent bot (A

---

[3] This was performed as an act of confirming with the decision made in the intellectual property case of US A&M Records, Inc. v. Napster, Inc., 239 F.3d 1004 (2001) and not an explicit attack against the Napster network, but the fact that the Napster network could so easily stop the network by just disconnecting its central server shows the inherent weakness of centralised architectures.
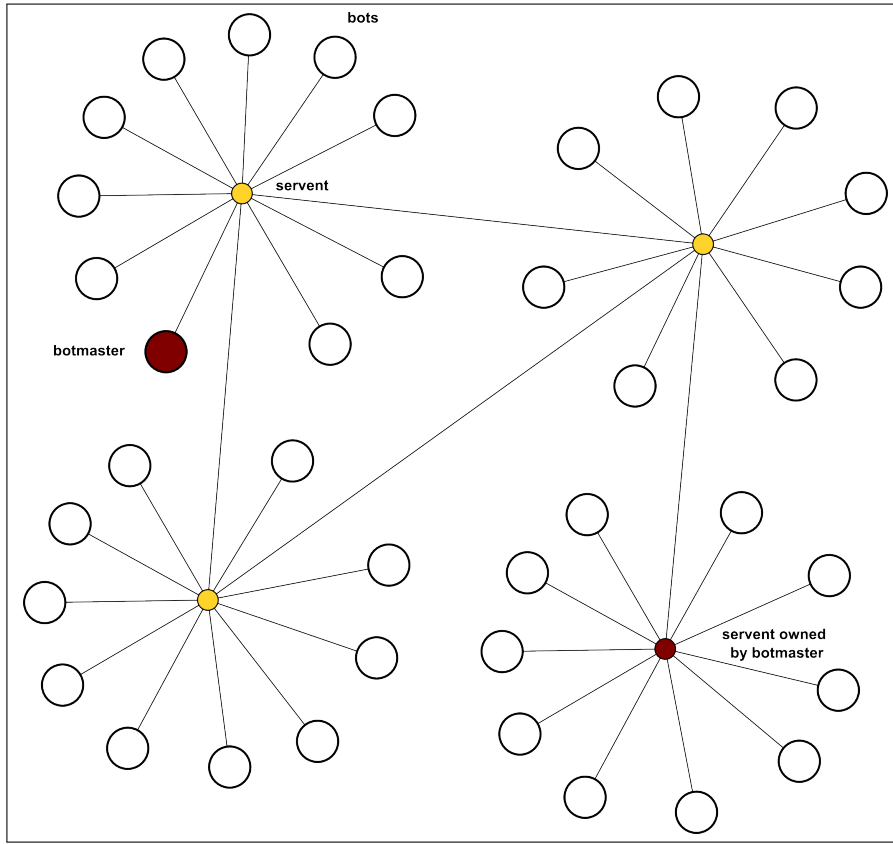
Fig. 2. An exemplary P2P network. The botmaster can be a regular peer, a servent/super peer/server (whatever the corresponding name in the underlying P2P network may be) or both. In contrast to other peers he injects new commands to be forwarded by others. It might make sense for the botmaster to own a few peers or servents, as this makes command distribution faster and increases his control over the network.

may not be aware of B's identity, for example, when B is compromised by an e-mail virus sent from A), A adds B into its peer list in the same way." [WSZ10]
– **Reinfection:** "If bot A reinfects bot B, bot B will then replace [a series of] randomly selected bots in its peer list with [...] bots from the peer list provided by A. Again, bots A and B will add each other into their respective peer lists if the other one is a servent bot."[WSZ10]

Both the advanced P2P botnet and the super botnet have their own P2P protocols for C&C. They implement push and pull C&C.[WWAZ09] When a bot receives a command it forwards it to all the peers in its list (push). If a bot cannot accept incoming connections (due to network misconfiguration, or a firewall), it actively polls other peers in its connection list from time to time to receive new commands (pull, cf. section 4.1 on page 4).

Lifetime Comparison against super botnet In the following, we will concentrate on the "the advanced P2P botnet" as described by [WSZ10]. It was shown by Wang et al. [WSZ10] that the super botnet suggested by Vogt et al. is not likely to become successful in a real world scenario:
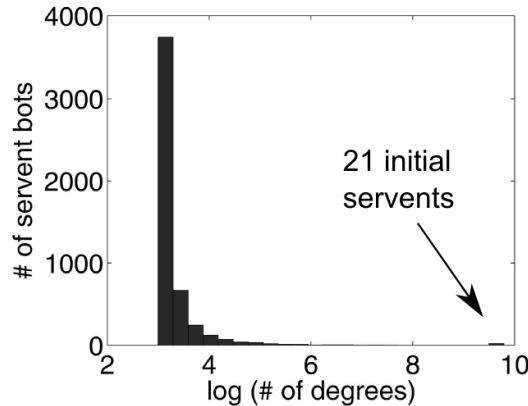

Fig. 3. Degree-distribution graph of servent bots assuming only new-infections and reinfections (derived work, original source: [WSZ10]). Simulation constraints: possible vulnerable hosts $n = 500,000$, stop of growth of botnet after $n = 20,000$, peer list size $M = 20$, 21 initial servent bots.

In contrast to what Vogt et al. assume, most of the compromised computers cannot act as servents (due to a firewall, NAT[4] or dynamic IP address) in the real world. Even though Vogt et al. demonstrated the robustness of their constructed botnet (cf. [VAJ07]), they rely on the assumption that enough reinfections will occur during the early lifetime of the botnet, namely the buildup-phase. According to [WSZ10] this is a false assumption as heavy reinfection-seeking during buildup will lead to easy detection of the botnet and a lot of wasted resources. Following this approach, Wang et al. showed that the super botnet algorithms for propagation — namely only "new infection" and "reinfection" — require over 200,000 infections events to create an evenly balanced botnet of only 20,000 vulnerable bots. As a result, when not enough reinfections occur, a scenario as depicted in figure 3 arises: Because the botnet stops growth after having infected 20,000 hosts, but there is such a

---

[4] Network Address Translation, occurs behind a router

Fig. 4. Degree-distribution graph of servent bots simulating peer list update propagation, also (Source: [WSZ10]). Most notably, the 21 initial servent bots cannot be spotted. Instead, a robust backbone for the P2P network has formed out of 1,000 servent bots v (around x-axis 6), $deg(v) \in [300; 500]$.

huge amount of vulnerable hosts (500,000), a reinfection event rarely ever happens. This means that servent bots only seldom exchange parts of their peer lists. As a consequence, the connection to servent bots is extremely unbalanced: In the network graph, 80% of servent bots have a degree[5] less than 30, while the initial 21 servents have degrees between 14,000 – 17,000: Most of the newly introduced servents have very few peers (both clients and servents) connected to them, essentially degrading the P2P botnet to a central network with 21 main servers. This is by no means an ideal P2P botnet.

When it is not possible to have enough reinfections, the "new infection" and "reinfection" propagation measures are obviously not enough.

Peer list updating Because of this problem, Wang et al. propose a new, third propagation method: Peer list updating. The idea behind this is that bots update their peer lists frequently. However, this imposes a severe security problem: An attacker capturing only one bot could soon reveal the identity of many servents in the network. Thus, a new command is introduced: Enforced peer list updating. As described in [WSZ10], it is possible for a botmaster to monitor his botnet, i.e. determine how many servents exist. After a sufficiently large time after construction phase, he can enforce a peer list update: All bots obtain a new peer list from a specified *sensor host*. This sensor host is equipped with the knowledge of all the servents in the network by the monitor-command issued by the botmaster. Upon query, the sensor host creates a peer list in the following way: It randomly chooses servent bots, composes an updated peer list out of them and sends it back to the querying bot. After each peer list update command, all bots will have "uniform and balanced connections."[WSZ10]

The network graph obtained is then similar to the one shown in figure 2 on page 9.

It is to be discussed at which point in time it makes sense for the botmaster to enforce a peer list update, for every update command bears the risk of discovery of parts of the network. This is further discussed in [WSZ10]. Simulations with an update after the first 1,000 infections show that this update strategy will result

---

[5] Degree is meant in the graph-theoretical sense here, i.e. counting the outgoing and incoming number of edges

in a degree distribution depicted in figure 4: The first 1,000 servents have many balanced connections ($deg(v) \in [300; 500]$), forming the robust backbone and connecting the hybrid P2P network tightly together. However, the remaining 4,000 servents (connected to the network after the update-command was run) have the known symptom of having degrees between 20-30, a situation well-known from the simulations of the super botnet (cf. figure 3). Thus, a forced peer list update from time to time seems necessary to make for a good P2P infrastructure.


Fig. 5. Hybrid network (Source: [WSZ10])

Further improvements The proposed hybrid network has several other advantages over common, existing P2P networks: It doesn't need bootstrapping, removing a single point of failure. Bootstrapping is avoided due to the three propagation measures (new infections, reinfections and forced peer list updates). An initial peer list is simply passed on to the newly infected zombie by the machine infecting it. Due to its fixed size peer list ($M = 20$ in the simulations for figures 3, 4), when an attacker gets access to a bot, it doesn't reveal whole (sub-)nets. Only machines with static IPs appear as peer bots, so-called *servents*. This way, it is guaranteed that the distributed peer lists are maximally deadlink-free. Data encryption in the hybrid P2P botnet has two functions: First, it makes it hard to sniff for patterns in internet traffic to detect a possible botnet. Second, the authenticity of the issued commands can be verified so that only commands signed by the botmaster are executed. The bot software can be shipped with a hard-coded public key from the botmaster.

Compared to a central-structure botnet (see figure 1), the hybrid P2P net (see figure 5) is only an extension of the original network: It is essentially equivalent to a central architecture P2P net. However, the amount of servers (in the form of servents) is greatly increased, as is the number of interconnections between them. The great number of servents is the primary reason why the hybrid P2P botnet is supposedly very hard to shut down.[WSZ10]

As tempting as such an advanced new botnet protocol may sound, it has to be acknowledged that the hybrid P2P botnet is a purely academical, theoretical idea that has never been tested in practice: "The network may not be as stable and robust as expected due to complex network conditions and defences."[WWAZ09]

6 Countermeasures against evil P2P botnets

What can an attacker do to de-arm a botnet at all? The ideal solution would be to shutdown all participants in the botnet (and only those). Realistically, this is almost never possible. It is also illegal in most countries to start a botnet, as it implies gaining access to another computer. This falls into the category of hacking. On the other hand, paradoxically, shutting down a botnet is likewise illegal, for it normally requires the attacker — be he benign or not — to do the same: Hacking into computers.

Common sub-tasks in fighting botnets are therefore, from an attacker's perspective:

Detecting the botnet at all, which can be very difficult for P2P parasite and leeching networks, as a separation between "bot" P2P and "real" P2P data streams has to be made. For example, in [SI07] it was shown that the bot Storm (therein referred to as "Peacomm") uses a fixed package size for 95,4% of its communication: 67 bytes (72,1% of traffic) and 60 bytes (23,3%). If one finds this atypical package size profile in a P2P network, it's likely to be infected by Storm.

Analysing the botnet: Finding the servers (if there are any) and participating zombies. It is also helpful to be able to estimate the size of the botnet. When using a technique like sybil attacks (see section 6.2), this gives an estimation of how many resources are needed, i.e. how much money will be needed to take down the botnet.

Preventing further spread of the botnet: Fixing the security exploit of the bot's injection vector. Most often, when a bot installs itself unnoticed from the owner (*injection*), it does so by using (*exploiting*) a security problem in the software. If the manufacturer fixes and updates this security hole, spread of the bot cannot continue on upgraded systems.

Disabling the bot(sub-)nets: Making clients loose inter-connection (in P2P networks), shutting down central servers (in central networks).

Infiltrating botnets to do non-malicious tasks, i.e. not to forward and not to execute incoming commands.

It must be noted that these tasks are of generic nature and might not be applicable to the botnet at hand. However, they represent a solid way of taking down a botnet.

While it is — in theory — relatively easy to shutdown a centralised architecture by determining the central servers and disabling those (e.g. through DoS attacks), P2P networks are arguably harder to disable, given they are properly protected. The legal implications of shutting down a botnet are also to be considered: When proven to be malicious, it is possible for the police to get a warrant to cut-off the servers physically from the net. This cannot practically be done for non-centralised P2P networks, since the bots are usually private property, with their owners unwillingly taking part in the botnet. Additionally they usually exist in so great masses that a physical take-down would be impossible to coordinate.

There are three well-known techniques to fight P2P botnets in particular.

6.1 Index poisoning

An attacker's idea: Polluting the file index The principal idea of index-poisoning is given by [WWAZ09]: "Originally, [the] index poisoning attack was introduced to prevent illegal distribution of copyrighted content in P2P networks. The main idea is to insert [a] massive number of bogus records into the index. If

a peer receives [a] bogus record, it could end up not being able to locate the file (nonexistent location), or downloading the wrong file."

It is easy to transfer this idea to the C&C of botnets. You only have to think that bots are not usually looking for files, but for commands. Once you know under which keys the botnet commands are stored in the index records, an attacker trying to shutdown the network can insert false commands under the same keys. If there is enough false information, chances to hit the real command issued by the botmaster are slim. The index gets "flooded" or poisoned by wrong commands. These can either be NOPs (no operations, meaning a command that does nothing), or may even help to uncover the botnet. Because of its principal of inserting faulty records into a system, index poisoning is also referred to as a "pollution attack"[LNR06].

Index-poising has been reported to have succeeded in fighting two recent P2P bots, namely Trojan.Peacomm and Storm. ([GSN$^+$07], [LNR06])

A botmaster's counter-defence: Authentication of commands It is an inherent property of any P2P network that any peer can and should distribute commands. However, how can one achieve that only the botmaster's commands are distributed and executed? In a centralised architecture, this is relatively easy to assure: Only the botmaster has the password to the servers, so only he can log onto them and issue new commands. However, with P2P networks an encryption mechanism for the commands must be implemented, such as public-key encryption (a freely available example implementation is [GNU11]). The botmaster signs (and possibly even encrypts) his commands with his private key. By supplying the bot clients with the botmaster's public key, they can authenticate the issuer of the command in question. As it is very difficult to calculate the private key based on the public key, bots can be assured that the issuer of the command is their botmaster. For this to practically work, file index records $< k, m >$, see section 4.3) could be enriched with a new column $H(m)$, where $H(m)$ is the signed hash value of m, i.e. the signature of the command.[WWAZ09] To verify m has been issued by the botmaster, a bot takes $H(m)$ and checks if it has been properly signed by using the hard-coded public signature.

6.2 Sybil attack

An attacker's idea: Inserting fake nodes A *sybil attack* on a botnet is a form of attack where large quantities of fake nodes (the so-called sybils) are inserted into the botnet. By infiltrating the botnet's peer list, the sybils try to interrupt the C&C between real bots. In contrast to the more static pollution attack, a sybil attack requires the attacker to have active hardware acting as the sybils and also to take part in the P2P network. Note however, that it is not necessary and even contra-indicative for the sybils to execute or distribute the commands sent over the P2P network's C&C structures. The idea of the sybil attack lies in the fact that there exist so many sybils in the network that issued commands get neither distributed nor executed. [DFNM]

Generally speaking, it is more expensive to run a sybil attack than an index-poisoning attack, as the former requires far more infrastructure at the disposal of the attacker.

A botmaster's counter-defence Botmasters could utilise known P2P techniques to fight the weakness against sybil attacks: According to [Dou02], it is common practice for a P2P system to distribute computational or storage tasks among remote peers to protect against the threat of hostile peers. For if a complete task is given to a hostile peer only, chances are the task is never executed and completely lost.

However, this risk is mitigated when the task is replicated to several peers, as it is unlikely for all peers to be hostile. In this way, a botnet could at least slightly compensate for an ever-growing mass of sybils in that its commands get executed anyway with a high probability (determined by the level of redundancy introduced).

This does not solve the problem of sybils interrupting communication in the first place. Countermeasures against this have been presented in section 5.3 on page 8.

Sybil attacks have been simulated to be a valid countermeasure against the Storm botnet.[DFNM]

6.3 Compromising Bootstrapping

Bootstrapping is a vulnerable point in any P2P botnet. When a hard-coded peer list is used (cf. section 5.2 for details), it is sufficient to take down all the peers in the bootstrapping table for the network to eventually shutdown: New bots simply can't find an initial peer to connect to. Botmasters have reacted to this by providing a Gnutella-like web-cache or updateable bootstrapping tables. As interrupting the bootstrapping process can so easily be done, it has been completely removed in the hybrid botnets (cf. section 5.3 on page 12).

7 Conclusions and Outlook

This paper gave an overview of the state-of-the-art in P2P botnets. While [WWAZ09] conclude that a shutdown of a P2P botnet is principally no harder than that of a central network when employing the index poisoning technique, our understanding and findings, as shown in this paper, are substantially different and backed-up by both reason and literature ([ZC06], [BH09], [SI07]). Additionally, new P2P protocols can be made resistant to index poisoning (shown in the same paper, [WWAZ09]), or file indices can be completely omitted, whereas the central nature of the servers remains a single point of failure in centralised architectures.

In fact, because they are harder to disable, more agile and dynamic in their life time, we believe that P2P botnets will play an ever-growing role in the future. It can also be cheaper to setup a P2P botnet, since there are no costs for central servers. Centralised architectures on the other hand have been seen to decline and will probably continue to do so. Also, some significant research for improvement in the existing P2P protocols has been conducted, removing many of the current weaknesses of P2P architectures.

Acknowledgement

References

ARZMT06. M. Abu Rajab, J. Zarfoss, F. Monrose, and A. Terzis. A multifaceted approach to understanding the botnet phenomenon. In *Proceedings of the 6th ACM SIGCOMM conference on Internet measurement*, pages 41–52. ACM, 2006.

Ben11. A. Benzin. The first cyberweapon? 2011.

BH09. A. Berger and M. Hefeeda. Exploiting sip for botnet communication. In *Secure Network Protocols, 2009. NPSec 2009. 5th IEEE Workshop on*, pages 31–36. IEEE, 2009.

Bor. R. Borgaonkar. An analysis of the asprox botnet. In *Emerging Security Information Systems and Technologies (SECURWARE), 2010 Fourth International Conference on*, pages 148–153. IEEE.

CJM05.   E. Cooke, F. Jahanian, and D. McPherson. The zombie roundup: Under-
            standing, detecting, and disrupting botnets. In *Proceedings of the USENIX
            SRUTI Workshop*, pages 39–44, 2005.
DD07.   D. Dittrich and S. Dietrich. Command and control structures in malware.
            Usenix magazine, 32(6), 2007.
DFNM.   C.R. Davis, J.M. Fernandez, S. Neville, and J. McHugh. Sybil attacks as a
            mitigation strategy against the storm botnet. In *Malicious and Unwanted
            Software, 2008. MALWARE 2008. 3rd International Conference on*, pages
            32–40. IEEE.
Dou02.   J. Douceur. The sybil attack. *Peer-to-peer Systems*, pages 251–260, 2002.
FMC.   N. Falliere, L.O. Murchu, and E. Chien. W32. stuxnet dossier. *Symantec
            Security Response*.
fSidIB11.   Bundesamt für Sicherheit in der Informationstechnik (BSI). Eröffnung des
            nationalen cyber-abwehrzentrums. `http://www.bmi.bund.de/SharedDocs/`
            `Termine/DE/2011/cyber.html?nn=373114`, June 2011. Accessed 2011/06/17.
GNU11.   GNU. Gnu privacy guard (gpg). `http://www.gnupg.org/`, June 2011. Ac-
            cessed 2011/07/10.
GSN+07.   J.B. Grizzard, V. Sharma, C. Nunnery, B.B.H. Kang, and D. Dagon. Peer-
            to-peer botnets: Overview and case study. In *Proceedings of the first confer-
            ence on First Workshop on Hot Topics in Understanding Botnets*, pages 1–1.
            USENIX Association, 2007.
Hru11.   Joel Hruska. Security admin, botmaster sentenced to four years
            in prison. `http://arstechnica.com/security/news/2009/03/`
            `security-admin-botmaster-sentenced-to-four-years-in-prison.ars`,
            June 2011. Accessed 2011/06/03.
Inc11.   Napster Inc. Napster. `http://www.napster.com`, June 2011. Accessed
            2011/06/01.
irc.   Irc - protocol defintion. `http://tools.ietf.org/html/rfc1459section-1`.
            Accessed 2011/06/05.
LJZ.   C. Li, W. Jiang, and X. Zou. Botnet: Survey and case study. In *Innovative
            Computing, Information and Control (ICICIC), 2009 Fourth International
            Conference on*, pages 1184–1187. IEEE.
LNR06.   J. Liang, N. Naoumov, and K.W. Ross. The index poisoning attack in p2p
            file sharing systems. In *IEEE INFOCOM*, volume 6. Citeseer, 2006.
New11.   C. T. News. Expert: Botnets no. 1 emerging internet threat. `http://www.cnn.`
            `com/2006/TECH/internet/01/31/furst/`, June 2011. Accessed 2011/07/14.
oC11.   University of California. Seti@home. `http://setiathome.berkeley.edu/`,
            June 2011. Accessed 2011/06/01.
SFS05.   D. Schoder, K. Fischbach, and C. Schmitt. Core concepts in peer-to-peer.
            Peer-to-peer computing: the evolution of a disruptive technology, page 1, 2005.
SI07.   M. Steggink and I. Idziejczak. Detection of peer-to-peer botnets. *University
            of Amsterdam, Netherlands*, 2007.
SK07.   R. Schoof and R. Koning. Detecting peer-to-peer botnets. *University of
            Amsterdam*, 2007.
SR06.   D. Stutzbach and R. Rejaie. Understanding churn in peer-to-peer networks. In
            Proceedings of the 6th ACM SIGCOMM conference on Internet measurement,
            pages 189–202. ACM, 2006.
VAJ07.   R. Vogt, J. Aycock, and M. Jacobson. Army of botnets. In *Network and
            Distributed System Security Symposium (NDSS)*. Citeseer, 2007.
Wik11.   Wikipedia. Napster. `http://en.wikipedia.org/w/index.php?title=`
            `Napster&oldid=432598619`, June 2011. Accessed 2011/06/14.
WSZ10.   Ping Wang, Sherri Sparks, and Cliff C. Zou. An advanced hybrid peer-to-peer
            botnet. *IEEE Transactions on Dependable and Secure Computing*, 7:113–127,
            2010.
WWAZ09.   P. Wang, L. Wu, B. Aslam, and C.C. Zou. A systematic study on peer-to-peer
            botnets. In *Computer Communications and Networks, 2009. ICCCN 2009.
            Proceedings of 18th Internatonal Conference on*, pages 1–8. IEEE, 2009.
ZC06.   C.C. Zou and R. Cunningham. Honeypot-aware advanced botnet construction
            and maintenance. 2006.